



CVPR 2025

# Cognitive AI for the Future: Multimodal Models and RAG in Vision Language Applications, from Training to Deployment

## Module 3: Optimize and Deploy the Multimodal RAG Pipeline

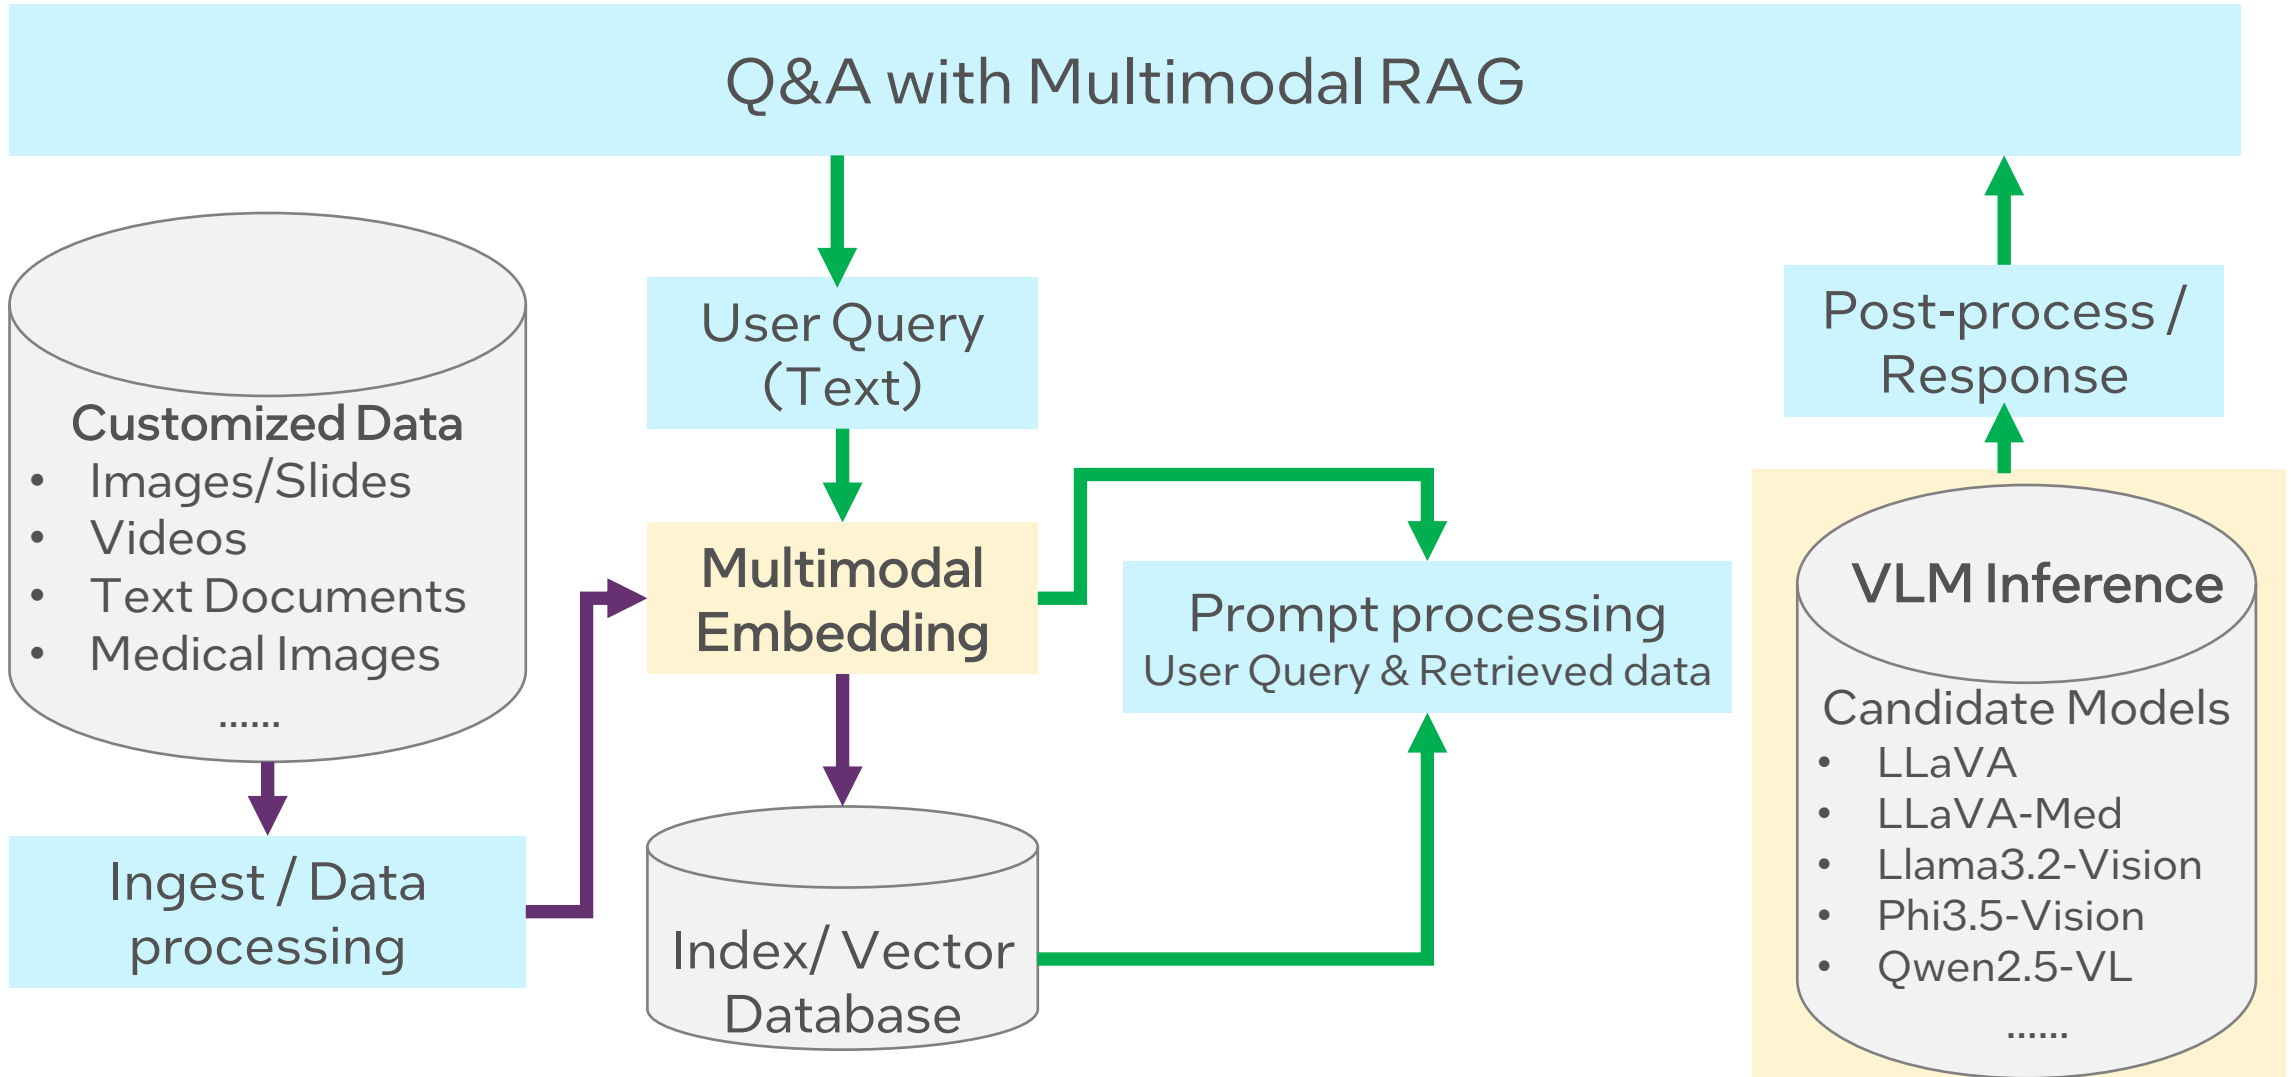
Speaker: Zhuo Wu

Job title: AI Software Evangelist

# Outline (30min)

- Overview the multimodal RAG pipeline 3m
- Convert & optimize multimodal embedding 7m
- RAG with LlamaIndex & OpenVINO 10m
- Multimodal RAG based video search 10m

# Multimodal RAG Pipeline



# Retrieval-Augmented Generation (RAG)

# LLM Ecosystem Integration



LangChain

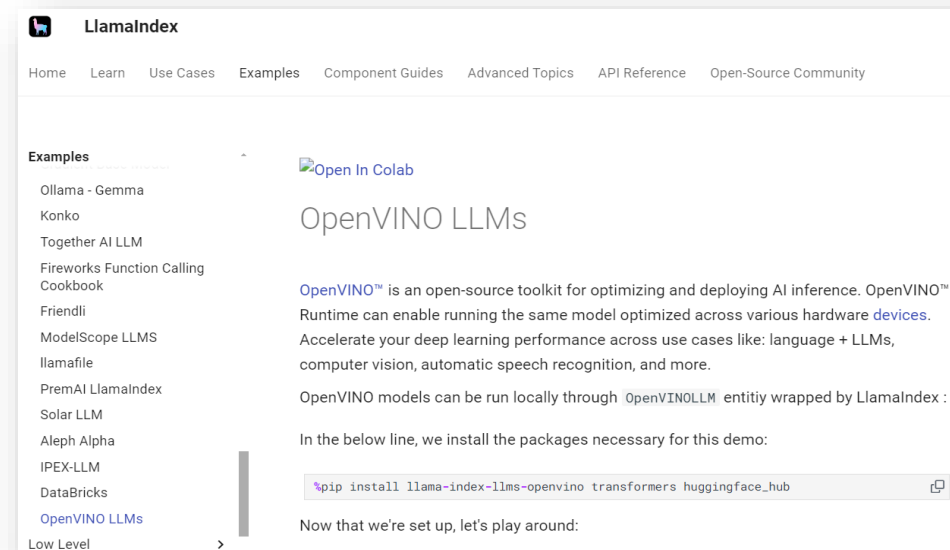
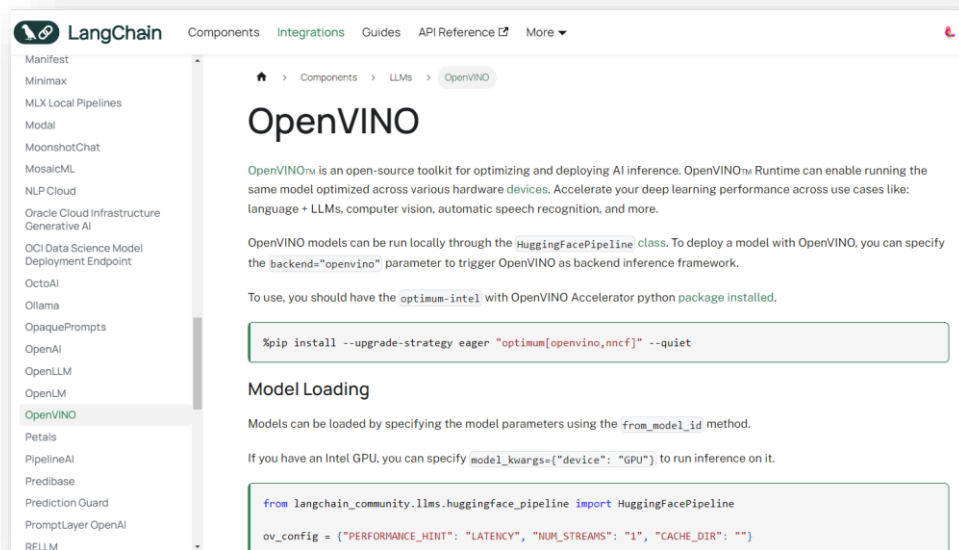
x



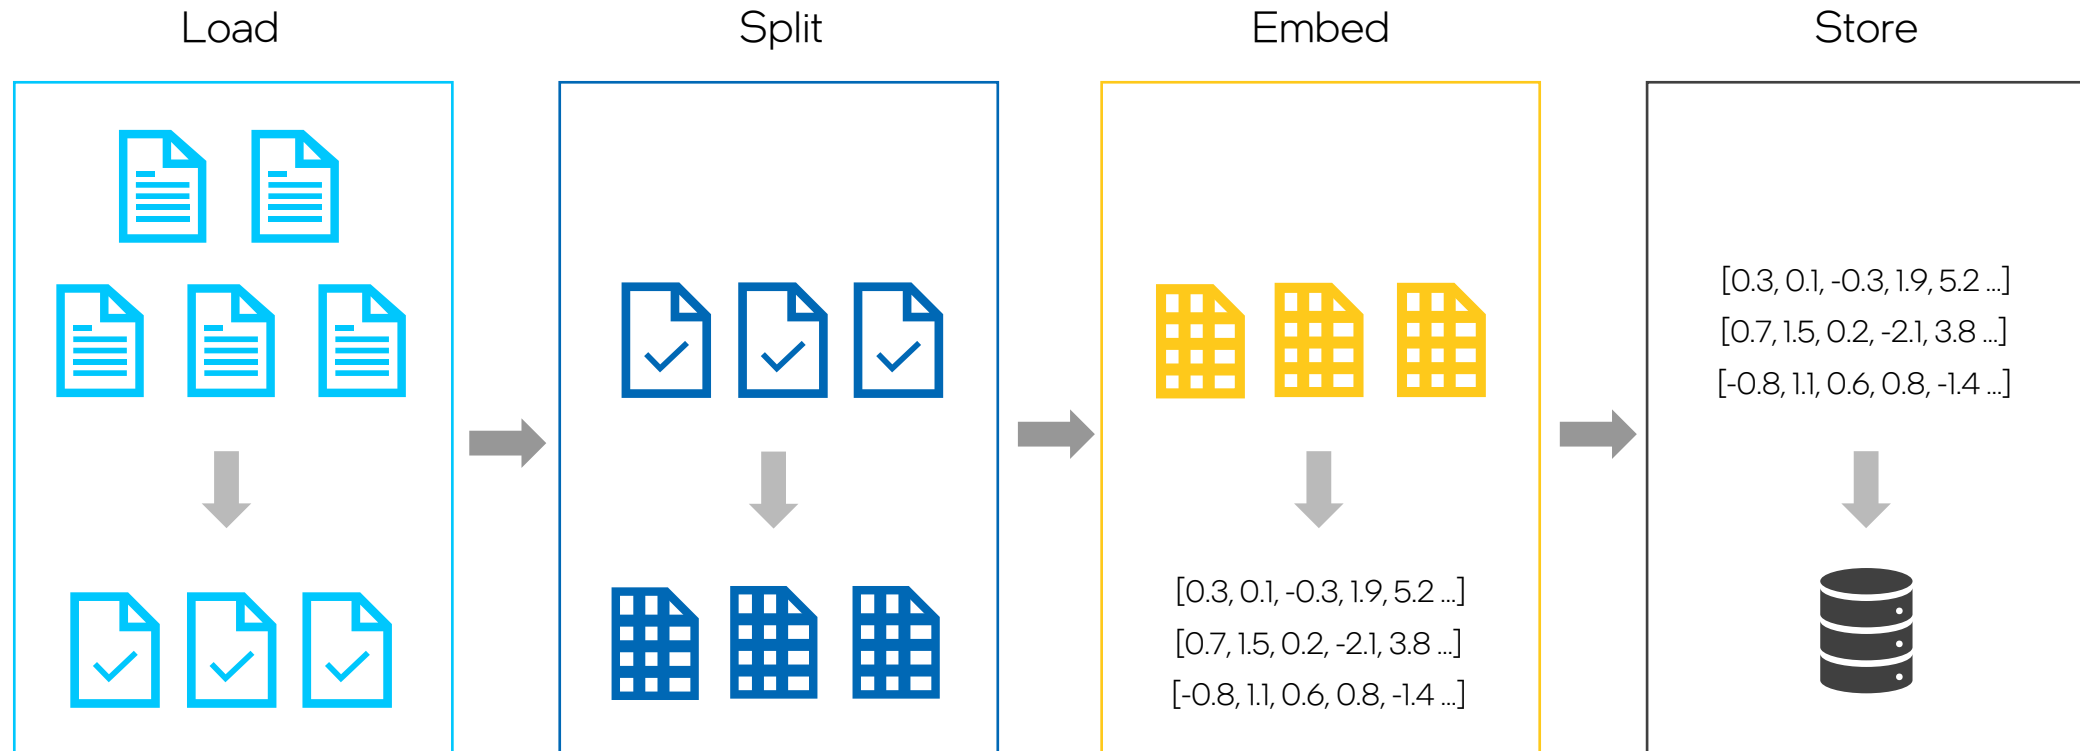
x



LlamaIndex

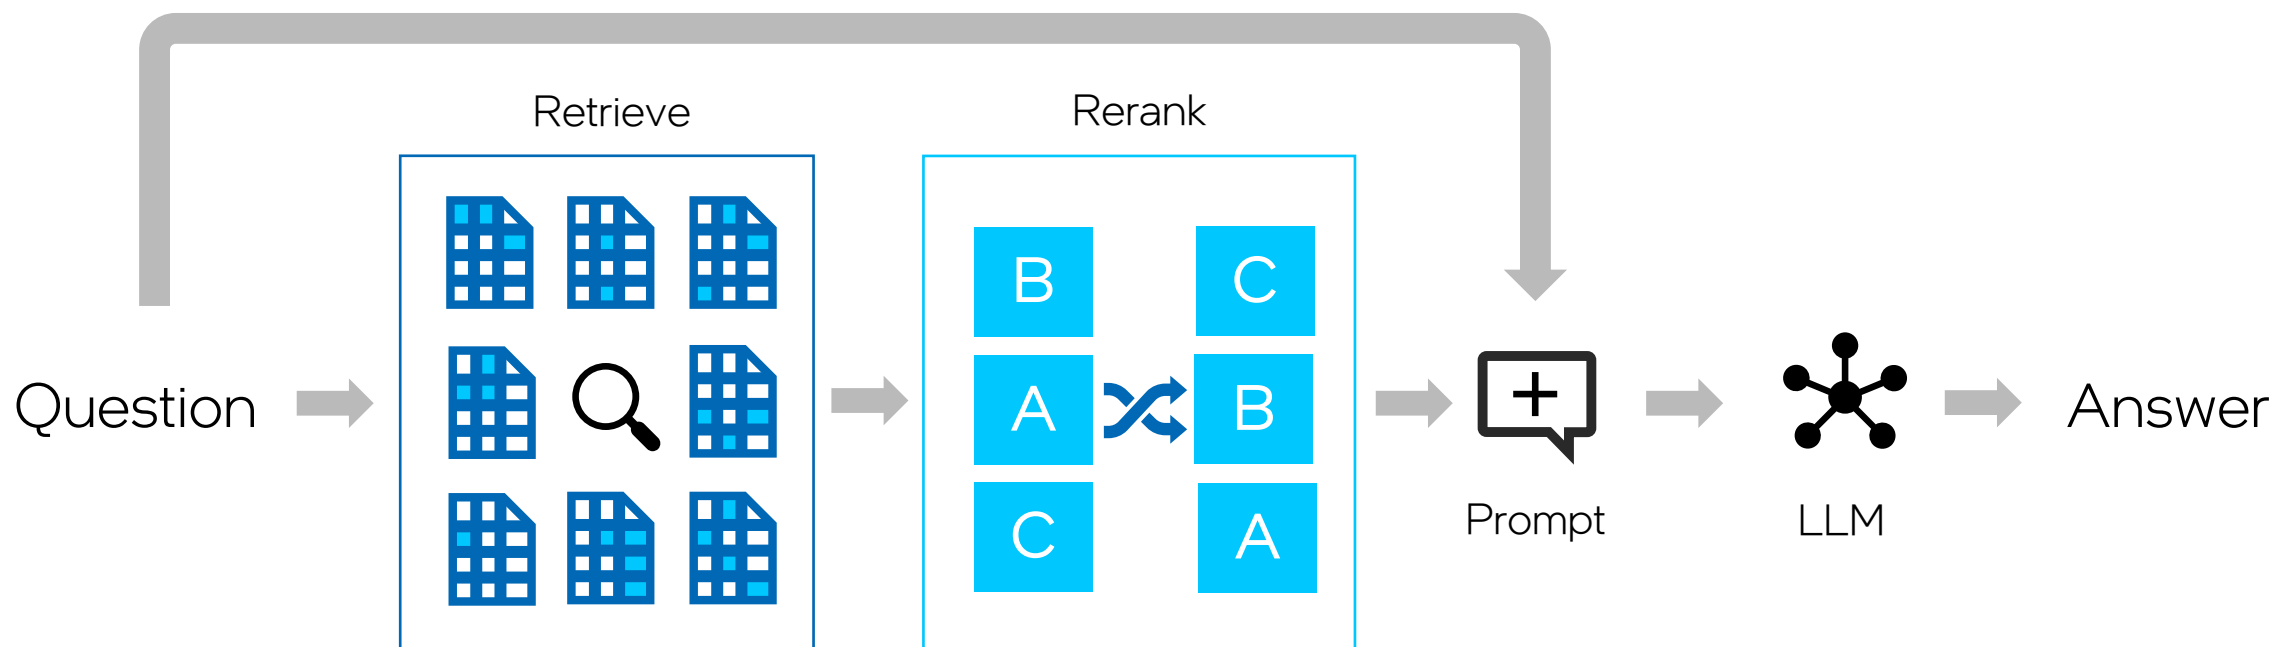


# Indexing





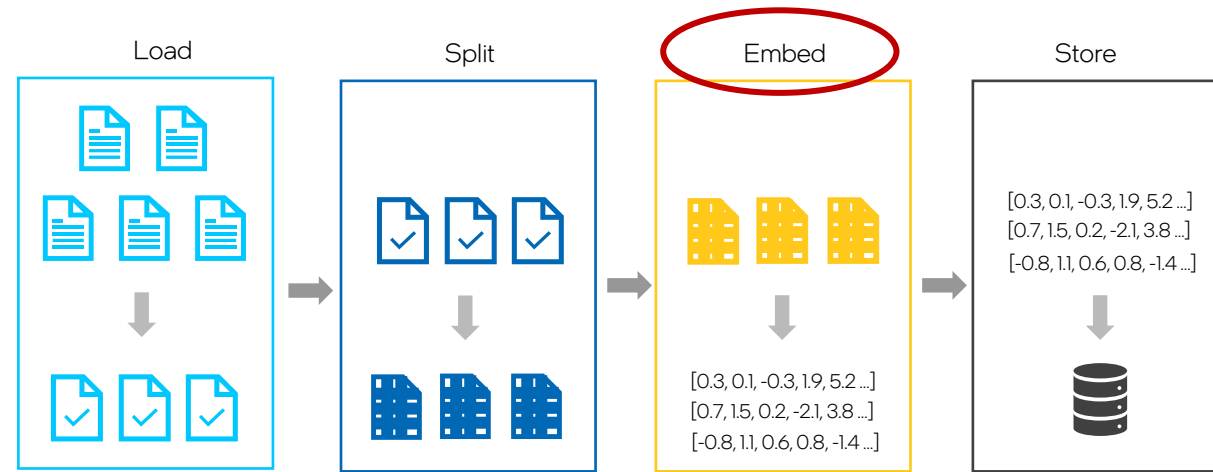
# Inference



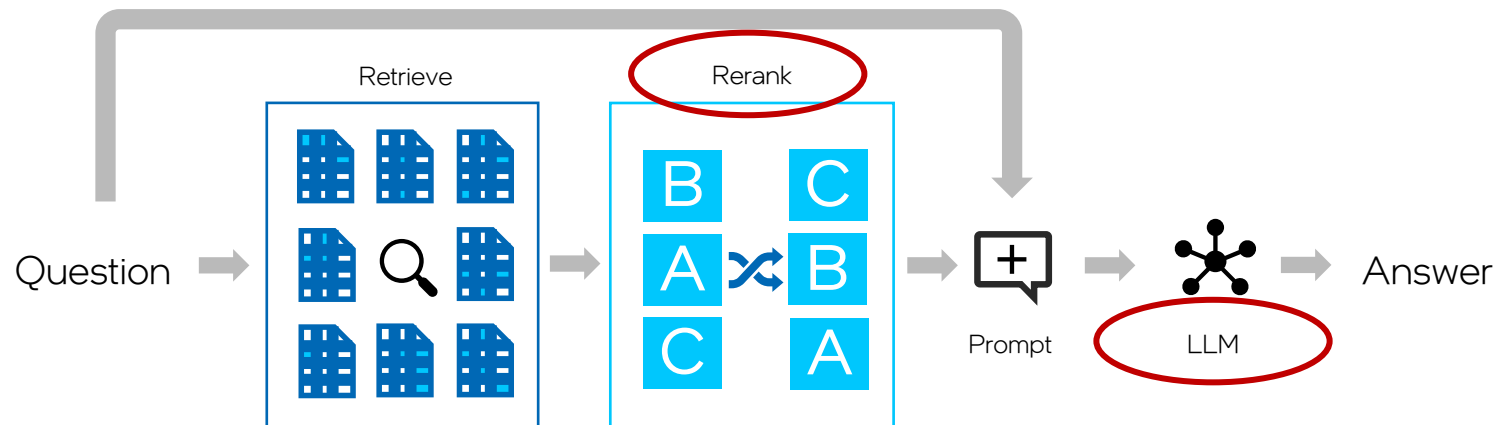
# Necessary models



## Indexing



## Inference





# Embeddings



```
pip install llama-index-embeddings-openvino
```

```
from llama_index.embeddings.huggingface_openvino import OpenVINOEmbedding
from optimum.intel import OVModelForFeatureExtraction
from transformers import AutoTokenizer

embedding_tokenizer = AutoTokenizer.from_pretrained(model_name)
embedding_tokenizer.save_pretrained(model_path)

embedding_model = OVModelForFeatureExtraction.from_pretrained(model_name, export=True,
                                                             compile=False)
embedding_model.save_pretrained(model_path)

embedding_model = OpenVINOEmbedding(str(model_path), device="CPU", embed_batch_size=1,
                                   model_kwargs={"dynamic_shapes": False})
```

# Reranker



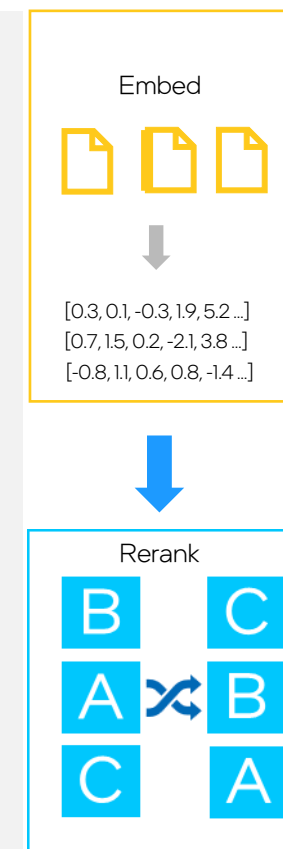
```
pip install llama-index-postprocessor-openvino-rerank
```

```
from optimum.intel import OVModelForSequenceClassification
from llama_index.postprocessor.openvino_rerank import OpenVINORerank
from transformers import AutoTokenizer

reranker_tokenizer = AutoTokenizer.from_pretrained(model_name)
reranker_tokenizer.save_pretrained(model_path)

reranker_model = OVModelForSequenceClassification.from_pretrained(model_name,
export=True, compile=False)
reranker_model.save_pretrained(model_path)

reranker_model = OpenVINORerank(model_id_or_path=str(model_path),
device="CPU", top_n=3)
```



# LLM



```
pip install llama-index-llms-openvino-genai
```

```
from llama_index.llms.openvino_genai import OpenVINOGenAILLM
from optimum.intel import OVModelForCausalLM, OVWeightQuantizationConfig
from transformers import AutoTokenizer

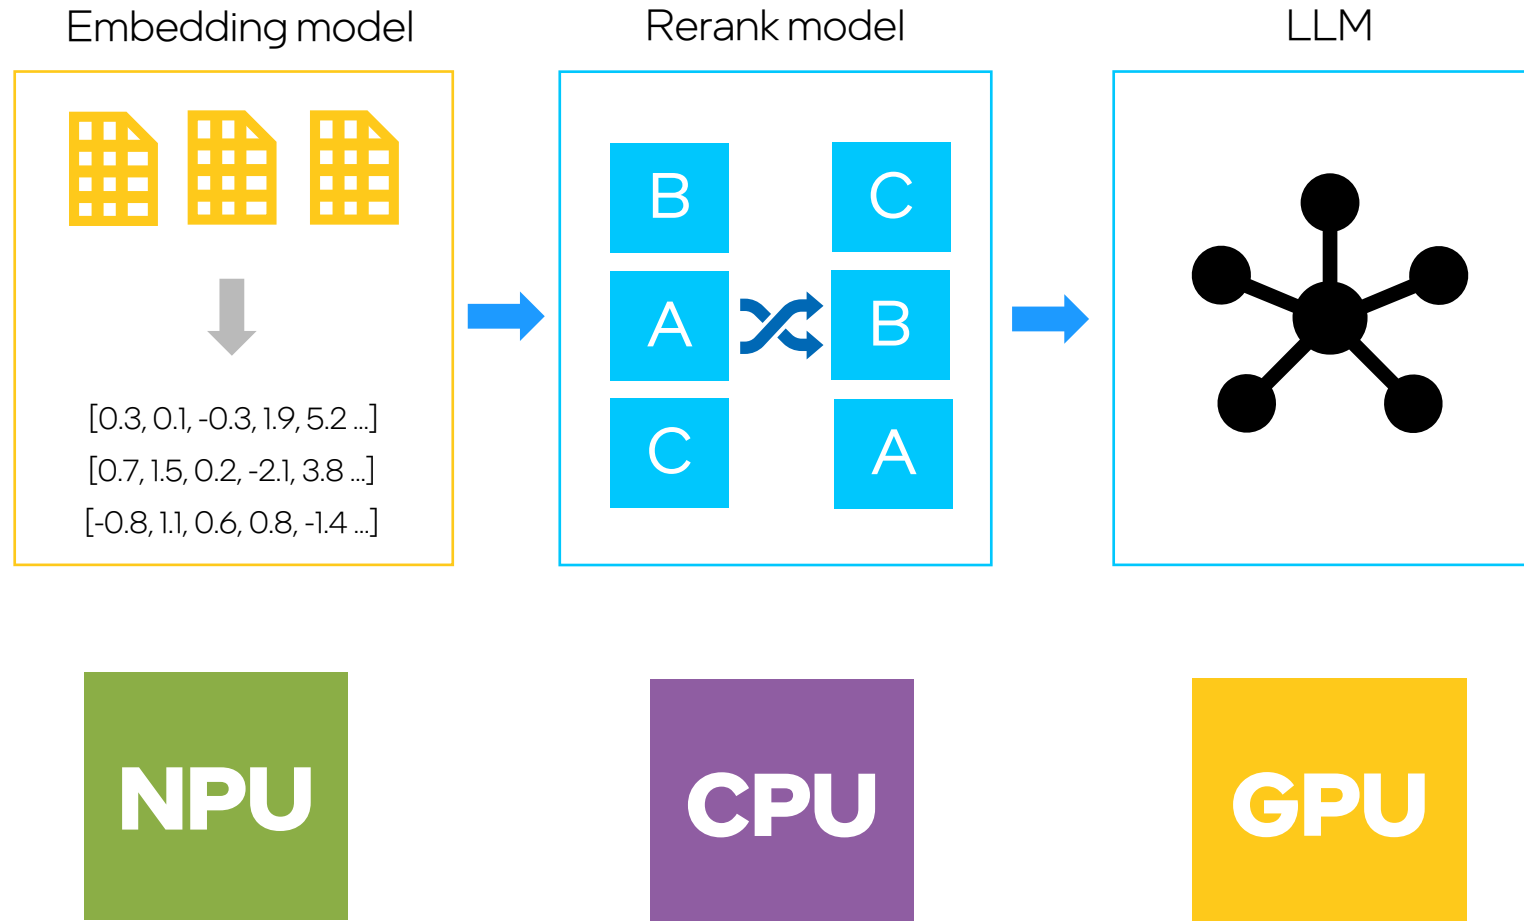
chat_tokenizer = AutoTokenizer.from_pretrained(model_name)
export_tokenizer(chat_tokenizer, model_path)

quant_config = OVWeightQuantizationConfig(bits=4, sym=False, ratio=0.8, quant_method="awq", group_size=128,
                                          dataset="wikitext2")
chat_model = OVModelForCausalLM.from_pretrained(model_name, export=True, compile=False, trust_remote_code=True,
                                                quantization_config=quant_config, library_name="transformers")
chat_model.save_pretrained(model_path)

ov_config = {"PERFORMANCE_HINT": "LATENCY", "CACHE_DIR": ""}
llm = OpenVINOGenAILLM(model_path=str(model_path), device="GPU", config=ov_config)

llm.config.max_new_tokens = 1024
llm.config.temperature = 0.7
llm.config.top_k = 50
llm.config.top_p = 0.95
```

# Deploy RAG Pipeline



# The chat engine



```
# limit chat history to 1024 tokens
memory = ChatMemoryBuffer.from_defaults(token_limit=2048)

documents = [Document(text=content, metadata={"file_name": file_path.name})]

# a splitter to divide document into chunks
splitter = LangchainNodeParser(RecursiveCharacterTextSplitter(chunk_size=500, chunk_overlap=100))

dim = ov_embedding._model.request.outputs[0].get_partial_shape()[2].get_length()
# a memory database to store chunks
faiss_index = faiss.IndexFlatL2(dim)
vector_store = FaissVectorStore(faiss_index=faiss_index)
storage_context = StorageContext.from_defaults(vector_store=vector_store)

# set embedding model
Settings.embed_model = ov_embedding
index = VectorStoreIndex.from_documents(documents, storage_context, transformations=[splitter])

# create a RAG pipeline
ov_chat_engine = index.as_chat_engine(llm=ov_llm, chat_mode=ChatMode.CONTEXT, memory=memory,
                                     system_prompt=chatbot_config["system_configuration"],
                                     node_postprocessors=[ov_reranker])
```

# Inference



```
chat_streamer = ov_chat_engine.stream_chat(question).response_gen  
  
answer = ""  
for partial_text in chat_streamer:  
    answer += partial_text
```

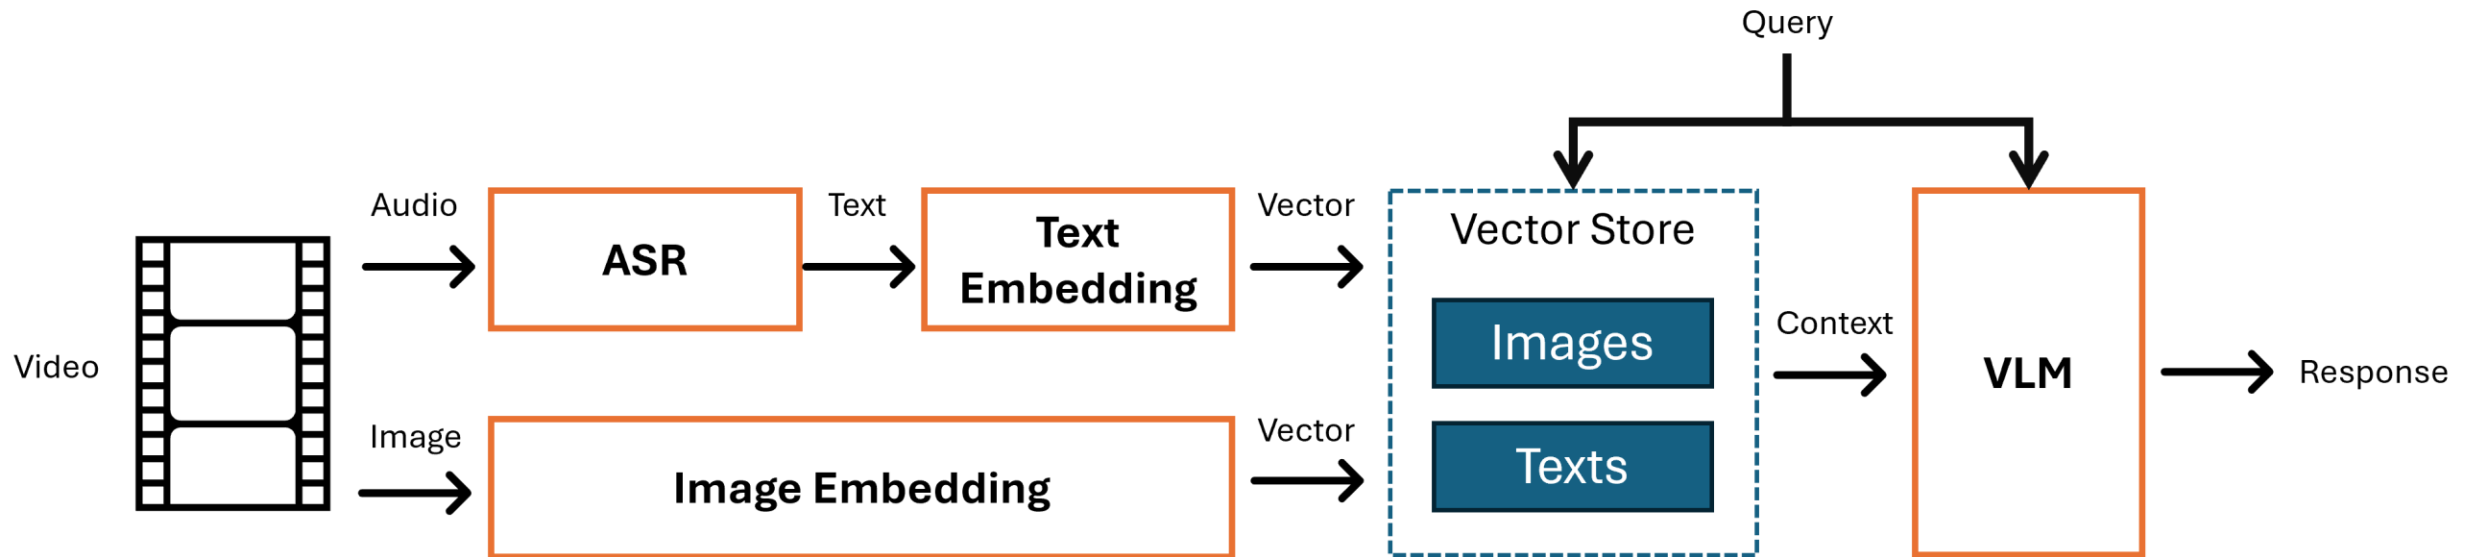


# Let's run the demo!

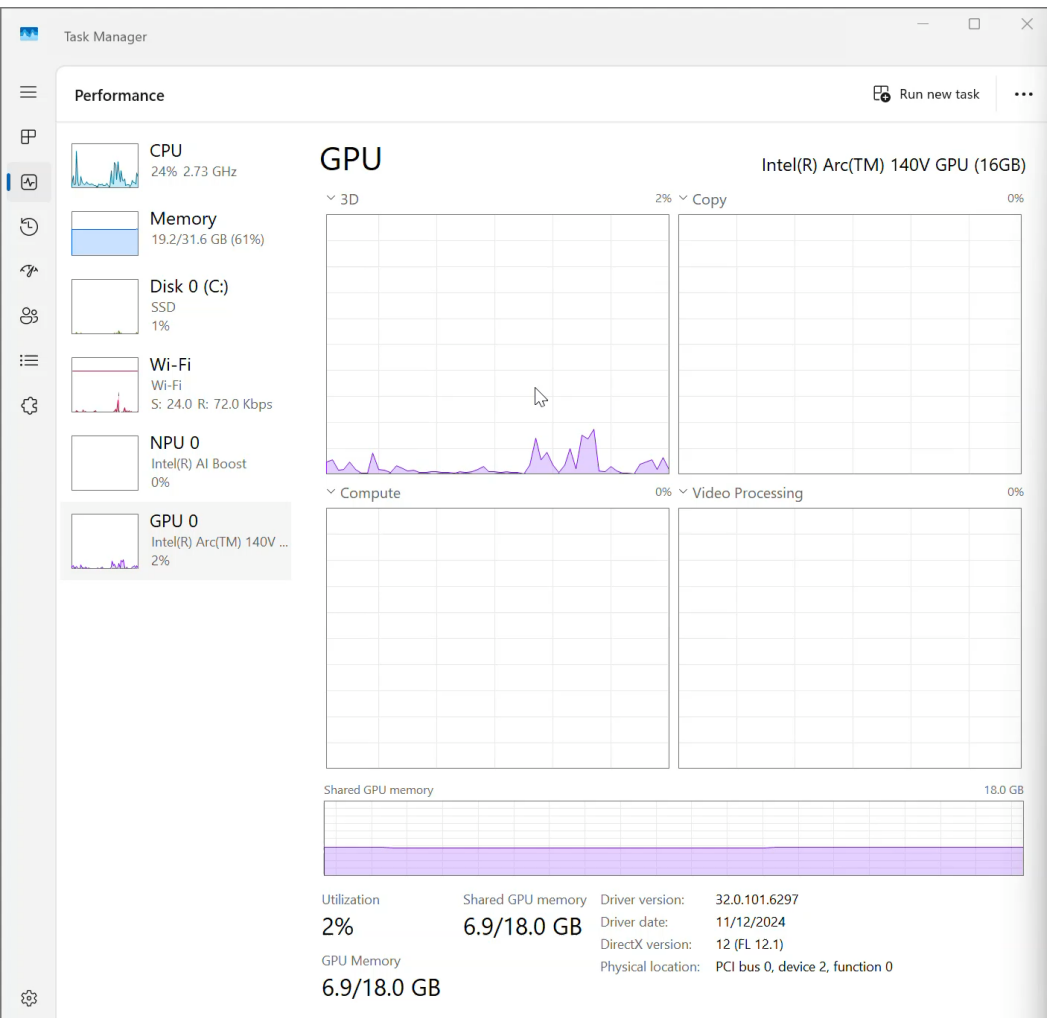


# Demo of LLM with RAG & Agent (AI assistant)

# Multimodal RAG & VLM with OpenVINO for Videos




# Convert, Optimize and Deploy VLM with OpenVINO



The image shows a web browser window with the URL 127.0.0.1:7860. The page title is "Phi-3.5-vision-instruct with OpenVINO". The main text says: "Try the [Phi-3.5-vision-instruct model](#) from Microsoft with OpenVINO. Upload an image and start chatting about it, or simply try one of the examples below. If you won't upload an image, you will receive an error." Below this is a "Chatbot" section with a large text input area. Underneath is an image input section showing a close-up of a human eye. Below the image is a text input field with the word "What" and a send button. At the bottom, there is an "Examples" section with two example prompts: "What is the text saying?" (with a small image of a cat) and "What does the chart display?" (with a small bar chart).

# QA over Video

Powered by OpenVINO



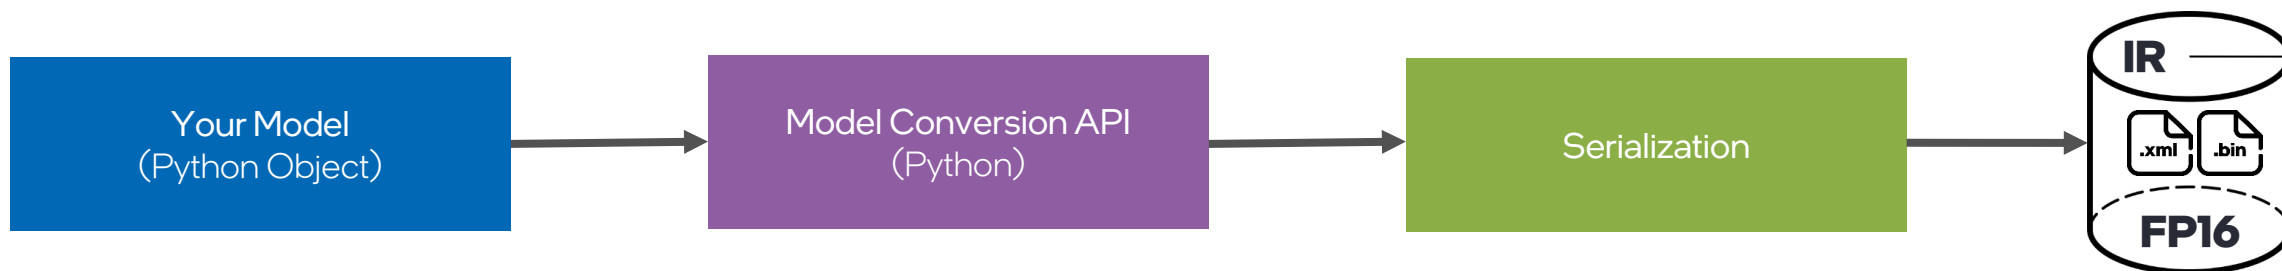
**Step 2: Build Vector Store**

Vector Store is Ready

**Step 3: Input Query**



# Convert & Optimize Multimodal Embedding Model (BridgeTower) with OpenVINO™ Model Converter



```
# Convert BridgeTower model to IR format and save
MODEL_NAME = "BridgeTower/bridgetower-large-itm-mlm-itc"
model = BridgeTowerForITC.from_pretrained(MODEL_NAME)
ov_model = ov.convert_model(model, example_input={**encoded_input})
ov.save_model(ov_model, "converted_model.xml")

# Convert BridgeTower text model to IR format and save
ov_textmodel_name = f"{models_dir}/custombridgetower_text_large_itc.xml"
text_model = BridgeTowerTextFeatureExtractor.from_pretrained(MODEL_NAME)
ov_textmodel = ov.convert_model(text_model, example_input={**text_encoding})
ov.save_model(ov_textmodel, ov_textmodel_name)
```




# Convert & Optimize Multimodal Embedding Model (BridgeTower) with OpenVINO™ Model Converter

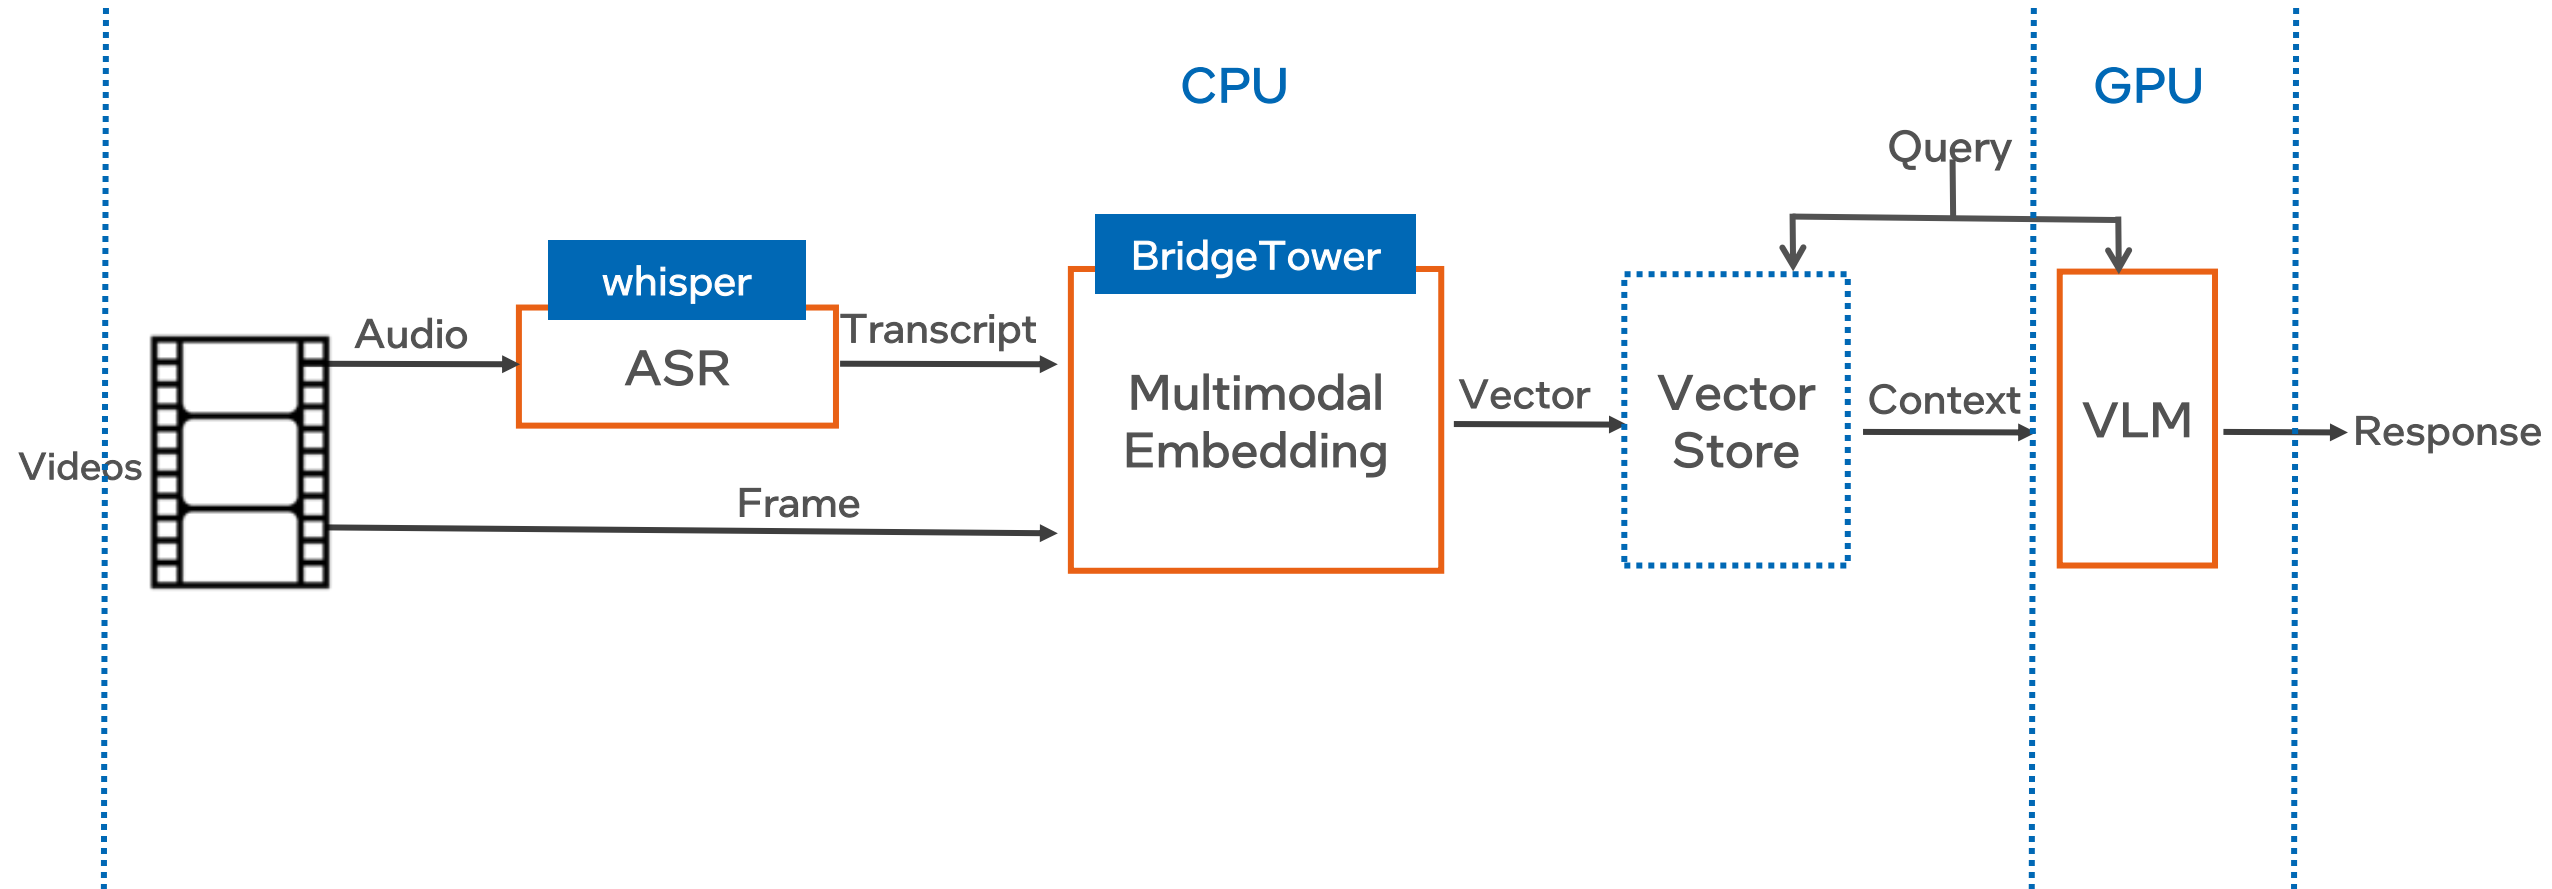
```
def restore_frompickle(filename:str)-> any:
    """ restore cookies from cache file"""
    with open(filename, 'rb') as handle:
        local_object = pickle.load(handle)
        return local_object

FILE_NAME=f"{models_dir}/input_example.pkl"
ENCODING_FILE_NAME =
f"{models_dir}/input_text_encoding.pkl"

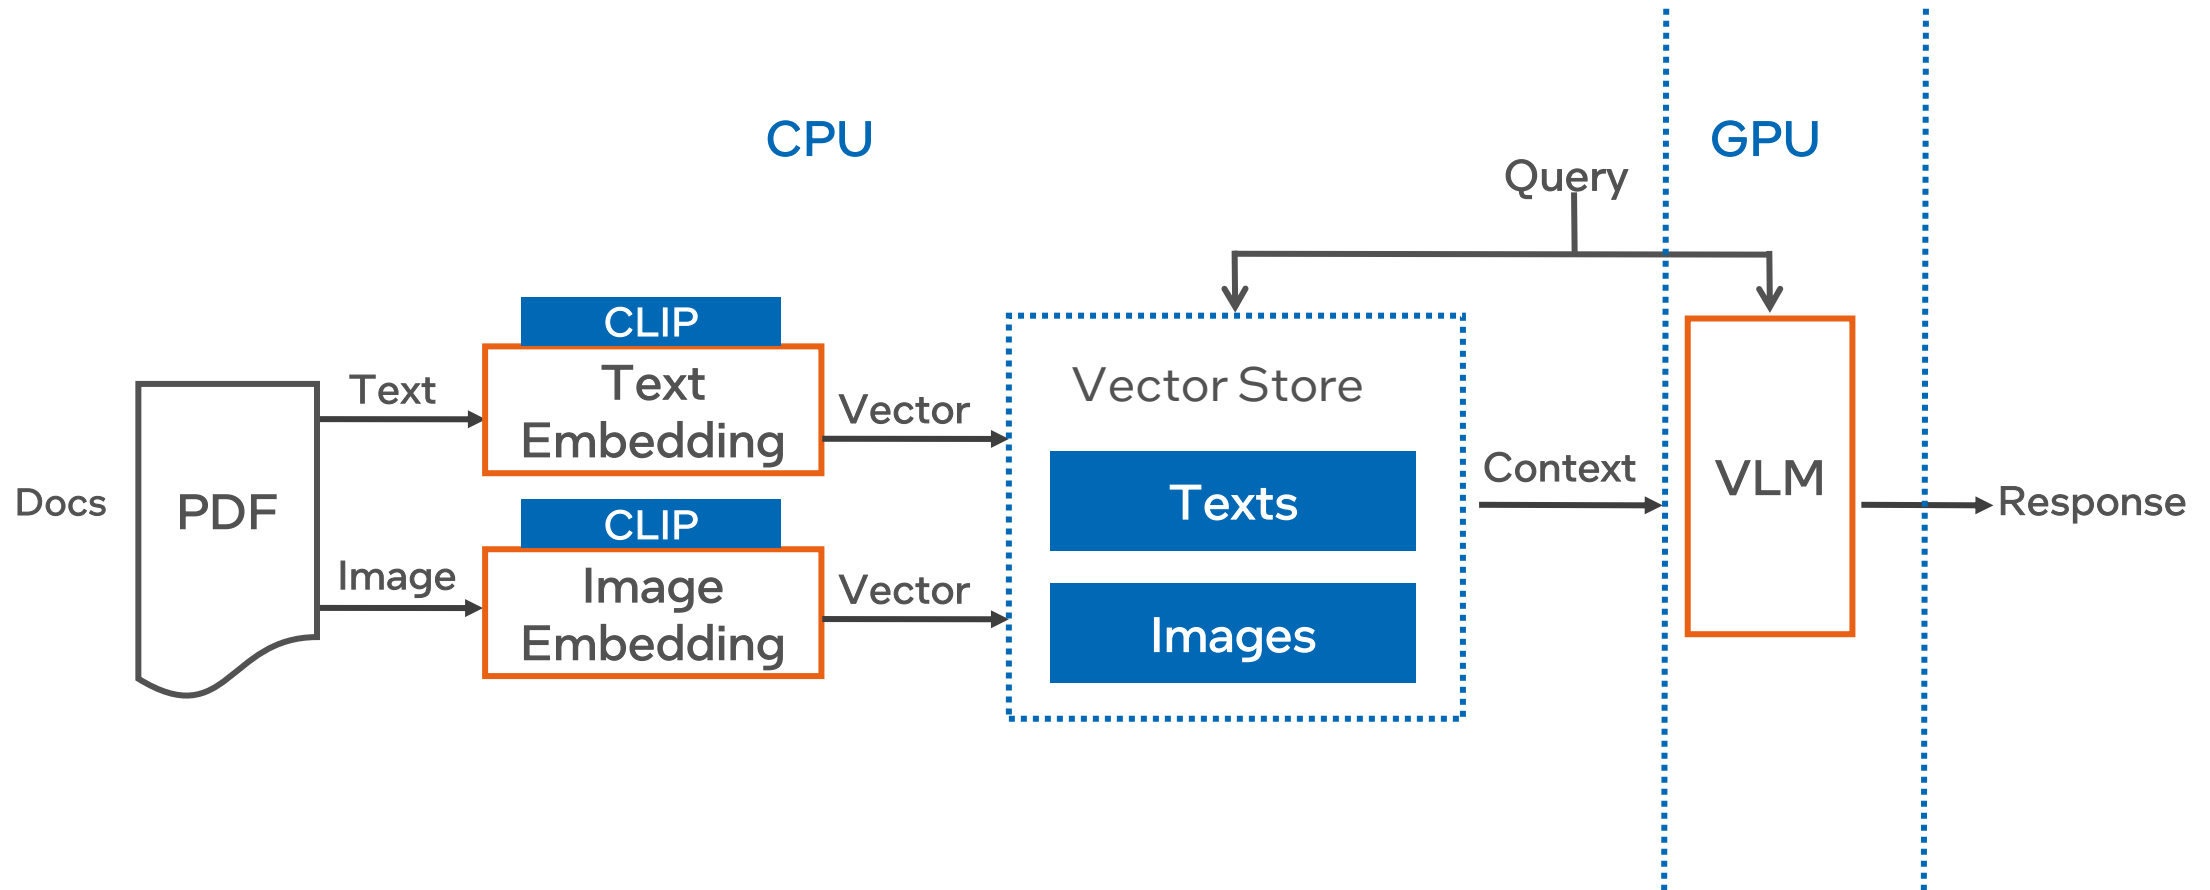
encoded_input = restore_frompickle(filename=FILE_NAME)
text_encoding =
restore_frompickle(filename=ENCODING_FILE_NAME)
```



# Multimodal RAG & VLM with OpenVINO for Videos



# Multimodal RAG & VLM with OpenVINO for Text/Image Documents



# Multimodal RAG Implementations with Langchain

```
from langchain_core.embeddings import Embeddings

class BridgeTowerEmbeddings(BaseModel, Embeddings):
    """ BridgeTower embedding model """

    def embed_documents(self, texts: List[str]) -> List[List[float]]:
        """Embed a list of documents using BridgeTower.

        Args:
            texts: The list of texts to embed.

        Returns:
            List of embeddings, one for each text.
        """
        ...

    def embed_image_text_pairs(self, texts: List[str], images: List[str], batch_size=2) -> List[List[float]]:
        """Embed a list of image-text pairs using BridgeTower.

        Args:
            texts: The list of texts to embed.
            images: The list of path-to-images to embed
            batch_size: the batch size to process, default to 2

        Returns:
            List of embeddings, one for each image-text pairs.
        """
        ...
```

# Multimodal RAG Implementations with Langchain

```
from langchain_community.vectorstores.lancedb import LanceDB

class MultimodalLanceDB(LanceDB):
    """`LanceDB` vector store to process multimodal data"""

    def add_text_image_pairs(
        self,
        texts: Iterable[str],
        image_paths: Iterable[str],
        metadatas: Optional[List[dict]] = None,
        ids: Optional[List[str]] = None,
        **kwargs: Any,
    ) -> List[str]:
        """Turn text-image pairs into embedding and add it to the database

        Args:
            texts: Iterable of strings to combine with corresponding images to add to the vectorstore.
            images: Iterable of path-to-images as strings.
            metadatas: Optional list of metadatas associated with the texts.
            ids: Optional list of ids to associate with the texts.

        Returns:
            List of ids of the added text-image pairs.
        """
        ...
```

# Multimodal RAG Implementations with Langchain

```
# initialize an BridgeTower embedder

embedder = BridgeTowerEmbeddings()

# ingest frame-transcription pairs
#   to lancedb vector store

_ = MultimodalLanceDB.from_text_image_pairs(
    texts=updated_vid1_trans+vid2_trans,
    image_paths=vid1_img_path+vid2_img_path,
    embedding=embedder,
    metadatas=vid1_metadata+vid2_metadata,
    connection=db,
    table_name=TBL_NAME,
    mode="overwrite",
)
```



# Multimodal RAG Implementations with Langchain

```
import lancedb

# Creating a LanceDB vector store
vectorstore = MultimodalLanceDB(
    uri=LANCEDB_HOST_FILE,
    embedding=embedder,
    table_name=TBL_NAME)

# creating a retriever for the vector store
retriever = vectorstore.as_retriever(
    search_type='similarity',
    search_kwargs={"k": 1}
)

# query the vector store
results = retriever.invoke("what dessert is included in the video?")
```

# VLM Pipeline

```
from transformers import AutoProcessor, TextStreamer
from optimum.intel.openvino import OVModelForVisualCausalLM

processor = AutoProcessor.from_pretrained(model_path, trust_remote_code=True)
ov_model = OVModelForVisualCausalLM.from_pretrained(model_path,
    device=device.value, trust_remote_code=True)

# text only input
conversation = [{"role": "user", "content": "What is the answer for 1+1? Explain it."}]
prompt = processor.tokenizer.apply_chat_template(conversation,
    tokenize=False, add_generation_prompt=True)

inputs = processor(text=prompt, images=None, return_tensors="pt")
print("Question:\nWhat is the answer for 1+1? Explain it.")
print("Answer:")
generate_ids = ov_model.generate(**inputs,
    max_new_tokens=50,
    streamer=TextStreamer(processor.tokenizer,
        skip_prompt=True, skip_special_tokens=True))

# text-image input
IMAGE_SPECIAL = "<|endoftext10|>"
AUDIO_SPECIAL = "<|endoftext11|>"
image = Image.open(image_path)
conversation = [{"role": "user", "content": f"{IMAGE_SPECIAL}What is unusual on this picture?"]}

prompt = processor.tokenizer.apply_chat_template(conversation, tokenize=False,
    add_generation_prompt=True)

inputs = processor(text=prompt, images=[image], return_tensors="pt")

generate_ids = ov_model.generate(**inputs, max_new_tokens=100,
    streamer=TextStreamer(processor.tokenizer,
        skip_prompt=True, skip_special_tokens=True))
```

This Space lets you run semantic search on a video.

YouTube URL

<https://www.youtube.com/watch?v=CgxDMH5dLao>

☐ Video

```
processing | 5.0s
```

```
processing | 5.0s
```

```

self.extract_images_and_embeds_updated(
File "C:\Users\yoyow\bridgetower-videosearch\videoSearchVinoRun.py", line 75, in timeit_wrapper
result = func(*args, **kwargs)
File "C:\Users\yoyow\bridgetower-videosearch\videoSearchVinoRun.py", line 447, in extract_images_and_embeds_updated
subtitles = get_transcript_vtt(video_id, path=output)
File "C:\Users\yoyow\bridgetower-videosearch\videoSearchVinoRun.py", line 75, in timeit_wrapper
result = func(*args, **kwargs)
File "C:\Users\yoyow\bridgetower-videosearch\videoSearchVinoRun.py", line 193, in get_transcript_vtt
raw = YouTubeTranscriptApi.get_transcript(video_id)
File "C:\Users\yoyow\openvino_env\lib\site-packages\youtube_transcript_api\api.py", line 306, in get_transcript
cls.list_transcripts(video_id, proxies, cookies)
File "C:\Users\yoyow\openvino_env\lib\site-packages\youtube_transcript_api\transcripts.py", line 134, in fetch
snippets = _TranscriptParser(preserve_formatting=preserve_formatting).parse(
File "C:\Users\yoyow\openvino_env\lib\site-packages\youtube_transcript_api\transcripts.py", line 474, in parse
for xml_element in ElementTree.fromstring(raw_data)
File "C:\Users\yoyow\openvino_env\lib\site-packages\defusedxml\common.py", line 127, in fromstring
return parser.close()
File "C:\Python310\lib\xml\etree\ElementTree.py", line 1722, in close
self._raiseerror(v)
File "C:\Python310\lib\xml\etree\ElementTree.py", line 1622, in _raiseerror
raise err
xml.etree.ElementTree.ParseError: no element found: line 1, column 0
clear_embedding_cache=True
Function get_video_id_from_url Took 0.0273 ms
video_url='https://www.youtube.com/watch?v=CgxDmH5dLao' video_id='CgxDmH5dLao' tmp_dir='./cache_tmp'
Getting video information for https://www.youtube.com/watch?v=CgxDmH5dLao
[youtube] Extracting URL: https://www.youtube.com/watch?v=CgxDmH5dLao
[youtube] CgxDmH5dLao: Downloading webpage
[youtube] CgxDmH5dLao: Downloading tv client config
[youtube] CgxDmH5dLao: Downloading tv player API JSON
[youtube] CgxDmH5dLao: Downloading ios player API JSON

```

### Examples

YouTube URL	Text query
<a href="https://www.youtube.com/watch?v=KCFYf4TJdN0">https://www.youtube.com/watch?v=KCFYf4TJdN0</a>	dessert
<a href="https://www.youtube.com/watch?v=KCFYf4TJdN0">https://www.youtube.com/watch?v=KCFYf4TJdN0</a>	man wearing a turkey

# Conclusion

# Building, Optimizing and Deploying Multimodal RAG Pipeline with OpenVINO™




Convert, optimize and deploy multimodal embedding models & VLMs in simple steps



Optimization of embedding models & VLMs brings smaller binary size & reduces memory footprint



Build RAG easily with OpenVINO's integration with LangChain, LlamaIndex, *etc.*

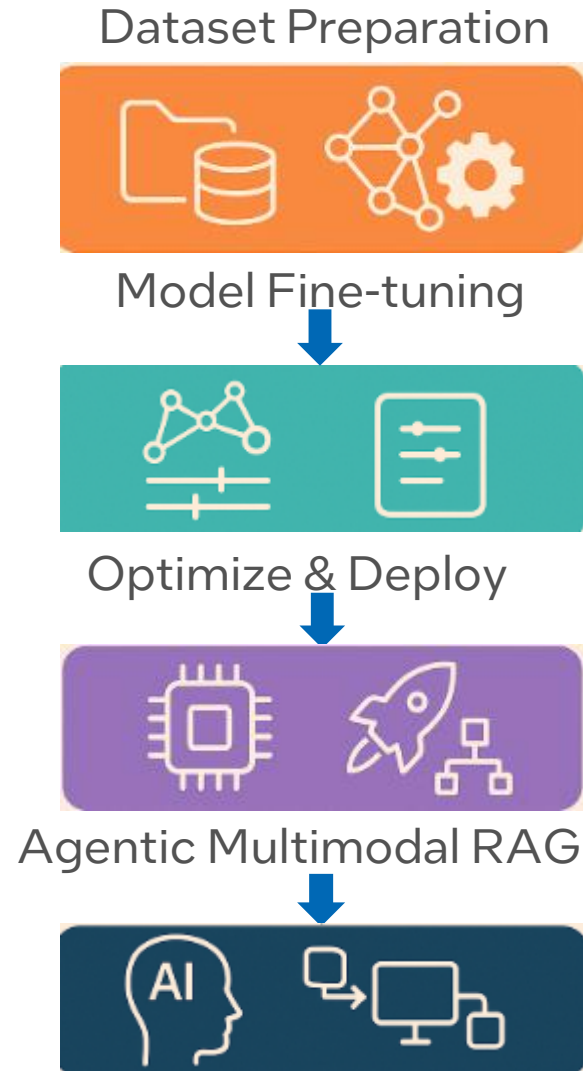


More accurate and intuitive retrieval by combining visual and textual understanding, empowering cognitive AI

# What's next?



# What You Will Learn



- Module 1: Tool for accelerating your dataset preparation locally
- Module 2: How to fine-tune the multimodal embedding model and VLM
- Module 3: How to optimize and deploy the multimodal RAG pipeline
- Module 4: How to build Agentic multimodal RAG